

# Concurrent and Parallel Programming in Python

<https://docs.python.org/3/library/concurrency.html#concurrent-execution>  
(<https://docs.python.org/3/library/concurrency.html#concurrent-execution>)

- There are two kinds of task
  - I/O bound: file accesses, network accesses, etc.
  - CPU bound: heavy computing tasks

|                   | Process                     | Thread                 |
|-------------------|-----------------------------|------------------------|
| Address space     | Separate address space      | Common address space   |
| Communication     | Inter-process communication | Memory synchronization |
| System resources  | More                        | fewer                  |
| Context switching | Slower                      | Faster                 |

## Race conditions

A race condition or race hazard is the behavior of an electronics, software, or other system where the output is dependent on the sequence or timing of other uncontrollable events. It becomes a bug when events do not happen in the order the programmer intended. (Wiki: [https://en.wikipedia.org/wiki/Race\\_condition](https://en.wikipedia.org/wiki/Race_condition))

### Example of a race condition

Threads #1 and #2 modify a global variable `_ref_count`, and the value of `_ref_count` is dependent on the executing sequence of the two threads if there is no proper synchronization (lock).

| Sequence #1    |                |            |
|----------------|----------------|------------|
| Thread #1      | Thread #2      | _ref_count |
| x = _ref_count |                | 0          |
| x = x + 1      |                | 0          |
| _ref_count = x |                | 1          |
|                | x = _ref_count | 1          |
|                | x = x + 1      | 1          |
|                | _ref_count = x | 2          |

| Sequence #2    |                |            |
|----------------|----------------|------------|
| Thread #1      | Thread #2      | _ref_count |
| x = _ref_count |                | 0          |
|                | x = _ref_count | 0          |
|                | x = x + 1      | 0          |
|                | _ref_count = x | 1          |
| x = x + 1      |                | 1          |
| _ref_count = x |                | 1          |

## Global Interpreter Lock (GIL)

- In Python, there are two kinds of tool for concurrent execution of codes.
  - In CPython, thread-based parallelism is suitable for tasks of I/O bound.
  - In CPython, process-based parallelism is suitable for tasks of CPU bound.

## In CPython, GIL is used to integrate libraries not thread-safe and true parallelism cannot always be achieved through multithreading.

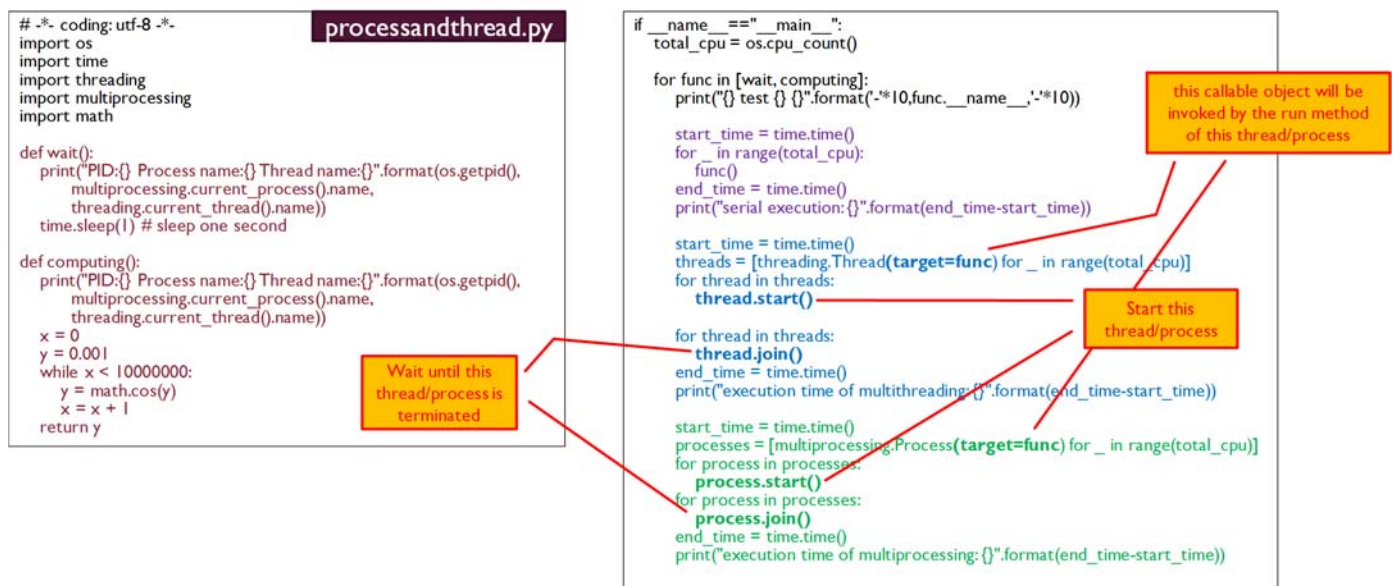
Python Implementations: <https://docs.python-guide.org/starting/which-python/> (<https://docs.python-guide.org/starting/which-python/>)

- CPython (<https://www.python.org/>) (<https://www.python.org/>)
- PyPy
- Jython (Java virtual machine)
- IronPython, PythonNet (.NET framework)

CPython requires the GIL: <https://wiki.python.org/moin/GlobalInterpreterLock>  
(<https://wiki.python.org/moin/GlobalInterpreterLock>)

Jython does not require the GIL: <http://www.jython.org/jythonbook/en/1.0/Concurrency.html>  
(<http://www.jython.org/jythonbook/en/1.0/Concurrency.html>)

## Processes vs Threads in Python



**multiprocessing.Process works normally when running Python interpreter but fails when running Jupyter notebook on Windows 10.**

- Creating a process is slower than creating a thread.
- In Python (CPython), multithreading is not suitable for tasks of CPU bound.

**Depending on the platform, multiprocessing supports three ways to start a process.**

`multiprocessing.set_start_method(method_name)`

- 'spawn' (Available on Unix and Windows. The default on Windows.)
  - 'fork' (The default on Unix)
  - 'forkserver' (Available on Unix platforms which support passing file descriptors over Unix pipes.)
- `set_start_method()` should not be used more than once in the program. Therefore, `set_start_method()` should be called in the `if __name__ == '__main__':` clause of the main module.

```
if __name__ == '__main__':  
    multiprocessing.set_start_method('spawn')
```

<https://docs.python.org/3/library/multiprocessing.html#contexts-and-start-methods>  
(<https://docs.python.org/3/library/multiprocessing.html#contexts-and-start-methods>)

In [2]:

```
1 # -*- coding: utf-8 -*-
2 import os
3 import time
4 import threading
5 import multiprocessing
6 import math
7
8 # an I/O-bound task
9 def wait():
10     print("PID:{} Process name:{} Thread name:{}".format(os.getpid(),
11         multiprocessing.current_process().name,
12         threading.current_thread().name))
13     time.sleep(1) # sleep one second
14
15 # a CPU-bound task
16 def computing():
17     print("PID:{} Process name:{} Thread name:{}".format(os.getpid(),
18         multiprocessing.current_process().name,
19         threading.current_thread().name))
20     x = 0
21     y = 0.001
22     while x < 10000000:
23         y = math.cos(y)
24         x = x + 1
25     return y
26
27 if __name__=="__main__":
28     total_cpu = os.cpu_count()
29
30     for func in [wait, computing]:
31         print("{} test {} {}".format('-'*10,func.__name__,'-'*10))
32
33         start_time = time.time()
34         for _ in range(total_cpu):
35             func()
36         end_time = time.time()
37         print("serial execution: {}".format(end_time-start_time))
38         #-----
39         start_time = time.time()
40         threads = [threading.Thread(target=func) for _ in range(total_cpu)]
41         for thread in threads:
42             thread.start()
43
44         for thread in threads:
45             thread.join()
46         end_time = time.time()
47         print("execution time of multithreading: {}".format(end_time-start_time))
48         #-----
49         start_time = time.time()
50         processes = [multiprocessing.Process(target=func) for _ in range(total_cpu)]
51         for process in processes:
52             process.start()
53         for process in processes:
54             process.join()
55         end_time = time.time()
56         print("execution time of multiprocessing: {}".format(end_time-start_time))
57
```

----- test wait -----

PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread

serial execution: 12.006043672561646

PID:6420 Process name:MainProcess Thread name:Thread-8PID:6420 Process name:  
MainProcess Thread name:Thread-9

PID:6420 Process name:MainProcess Thread name:Thread-10  
PID:6420 Process name:MainProcess Thread name:Thread-11  
PID:6420 Process name:MainProcess Thread name:Thread-12  
PID:6420 Process name:MainProcess Thread name:Thread-13  
PID:6420 Process name:MainProcess Thread name:Thread-14  
PID:6420 Process name:MainProcess Thread name:Thread-15  
PID:6420 Process name:MainProcess Thread name:Thread-16  
PID:6420 Process name:MainProcess Thread name:Thread-17  
PID:6420 Process name:MainProcess Thread name:Thread-18  
PID:6420 Process name:MainProcess Thread name:Thread-19

execution time of multithreading: 1.068385124206543

execution time of multiprocessing: 0.1514904499053955

----- test computing -----

PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread  
PID:6420 Process name:MainProcess Thread name:MainThread

serial execution: 13.88290786743164

PID:6420 Process name:MainProcess Thread name:Thread-20PID:6420 Process nam  
e:MainProcess Thread name:Thread-21

PID:6420 Process name:MainProcess Thread name:Thread-22  
PID:6420 Process name:MainProcess Thread name:Thread-23  
PID:6420 Process name:MainProcess Thread name:Thread-24  
PID:6420 Process name:MainProcess Thread name:Thread-25  
PID:6420 Process name:MainProcess Thread name:Thread-26  
PID:6420 Process name:MainProcess Thread name:Thread-27  
PID:6420 Process name:MainProcess Thread name:Thread-28  
PID:6420 Process name:MainProcess Thread name:Thread-29  
PID:6420 Process name:MainProcess Thread name:Thread-30  
PID:6420 Process name:MainProcess Thread name:Thread-31

execution time of multithreading: 14.21201229095459

execution time of multiprocessing: 0.1376032829284668

---

## threading.Thread: create a thread object

```
import threading
athread = threading.Thread(target=f,args=(p1,...),kwargs={k1:v1,...})
```

where args and kwargs are positional and keyword parameters for the function f.

- The main thread cannot terminate unless all non-daemon threads are terminated.
- Daemon threads are created by setting the daemon parameter True as `threading.Thread(daemon=True)`. The daemon thread can be kept when the entire program exits.

---

## Thread Methods

- Call `start()` to start a thread.
- Call `join(timeout)` to wait at most timeout seconds for the completion of a thread.
- Call `is_alive()` to test if the thread is alive.

Reference <https://docs.python.org/2/library/threading.html#thread-objects>  
(<https://docs.python.org/2/library/threading.html#thread-objects>)

## An example of a race condition

In [9]:

```
1 import threading
2
3 x = 0
4 def func(v):
5     global x
6     for i in range(20000000):
7         x = v
8         if x != x:
9             print('thread: {} v: {} x: {}'.format(threading.current_thread().name,v,x))
10
11 t1 = threading.Thread(target=func,args=(1,))
12 t2 = threading.Thread(target=func,args=(2,))
13 print('some messages will be shown if x!=x')
14 t1.start()
15 t2.start()
16
17 t1.join()
18 t2.join()
19
20 print('completed')
```

some messages will be shown if x!=x

```
thread: Thread-22 v: 1 x: 2
thread: Thread-23 v: 2 x: 1
thread: Thread-23 v: 2 x: 1
thread: Thread-23 v: 2 x: 1
thread: Thread-22 v: 1 x: 2
thread: Thread-23 v: 2 x: 1
thread: Thread-22 v: 1 x: 2
thread: Thread-22 v: 1 x: 2
thread: Thread-22 v: 1 x: 2
thread: Thread-23 v: 2 x: 1
thread: Thread-22 v: 1 x: 2
thread: Thread-22 v: 1 x: 2
thread: Thread-22 v: 1 x: 2
thread: Thread-22 v: 1 x: 2
thread: Thread-23 v: 2 x: 1
thread: Thread-22 v: 1 x: 2
thread: Thread-23 v: 2 x: 1
thread: Thread-23 v: 2 x: 1
thread: Thread-23 v: 2 x: 1
thread: Thread-22 v: 1 x: 2
thread: Thread-22 v: 1 x: 2
thread: Thread-22 v: 1 x: 2
thread: Thread-23 v: 2 x: 1
thread: Thread-22 v: 1 x: 2
thread: Thread-23 v: 2 x: 1
thread: Thread-22 v: 1 x: 2
completed
```

## Objects used to synchronize thread objects.

- **Lock/RLock:** A lock can be in either “locked” or “unlocked”.
  - It is created in the unlocked state.

- `acquire()` can change the state from unlocked to locked. When the state is locked, `acquire()` blocks until a call to `release()` in another thread, then `acquire()` resets it to locked and returns.
- `release()` changes the state to unlocked and returns immediately. If an attempt is made to release an unlocked lock, a `ThreadError` will be raised.
- **Condition Objects**
- **Semaphore Objects:** A semaphore has an internal counter which is decremented by each `acquire()` call and incremented by each `release()` call. When `acquire()` finds that this counter is zero, it blocks, and waits until some other thread calls `release()` of this semaphore.
- **Event Objects**

Lock, RLock, Condition, and Semaphore objects can be used as with statement context managers.

```
lock_object = threading.Lock()

with lock_object:
    # acquire() will be called automatically after entering this block
    ...
    # release() will be called automatically before leaving this block
```

## An example of using `threading.Lock` to synchronize thread objects.

In [1]:

```
1 import threading
2
3 x = 0
4 def func(v, lock):
5     global x
6     for i in range(2000000):
7         with lock:
8             x = v
9             if x != x:
10                 print('thread: {} v: {} x: {}'.format(threading.current_thread().name,
11 #             Lock.acquire()
12 #             x = v
13 #             if x != x:
14 #                 print('thread: {} v: {} x: {}'.format(threading.current_thread().name, v, x,
15 #             Lock.release()
16
17 lock_object = threading.Lock()
18
19 t1 = threading.Thread(target=func, args=(1, lock_object))
20 t2 = threading.Thread(target=func, args=(2, lock_object))
21
22 print('some messages can be shown if x!=x')
23
24 t1.start()
25 t2.start()
26
27 t1.join()
28 t2.join()
29
30 print('completed')
```


some messages can be shown if `x!=x`  
completed



---

## Deadlocks

A deadlock is a state in which some threads/processes wait for each other.



In [12]:

```
1 import threading
2 import time
3 def func1():
4     global t2
5     time.sleep(1)
6     print('{} waits for {}'.format(threading.current_thread().name, t2.name))
7     t2.join()
8     print('{} completes'.format(threading.current_thread().name))
9
10 def func2():
11     global t1
12     time.sleep(1)
13     print('{} waits for {}'.format(threading.current_thread().name, t1.name))
14     t1.join()
15     print('{} completes'.format(threading.current_thread().name))
16
17 t1 = threading.Thread(target=func1)
18 t2 = threading.Thread(target=func2)
19
20 t1.start()
21 t2.start()
22
23
24 print('you cannot see the message "thread-xxx completes"')
```

you cannot see the message "thread-xxx completes"

Thread-29 waits for Thread-28

Thread-28 waits for Thread-29

---

## Timer Objects

Example 1. timer\_func will be executed after 5 seconds.

In [11]:

```
1 import threading
2 import time
3 from datetime import datetime
4
5 def timer_func(n):
6     print('thread starts:{}'.format(datetime.now()))
7     for _ in range(n):
8         print('{}hello'.format(threading.current_thread().name))
9
10 t = threading.Timer(5, timer_func, (10,))
11 print('timer activated:{}'.format(datetime.now()))
12 t.start()
13 time.sleep(10)
14 print(t.is_alive())
15 t.join()
```

```
timer activated:2018-11-09 13:04:01.748624
thread starts:2018-11-09 13:04:06.757735
Thread-17,hello
Thread-17,hello
Thread-17,hello
Thread-17,hello
Thread-17,hello
Thread-17,hello
Thread-17,hello
Thread-17,hello
Thread-17,hello
Thread-17,hello
Thread-17,hello
Thread-17,hello
False
```

Example 2. The Timer object can be used to call a function for every n seconds.

In [1]:

```
1 import threading
2 import time
3 from datetime import datetime, timedelta
4 repeat = 5
5 def timer_func2(n):
6     global repeat
7     if repeat > 0:
8         t = threading.Timer(n, timer_func2,(n,))
9         t.start()
10        print('thread {} will start at {}'.format(t.name,datetime.now()+timedelta(seconds=n)))
11        repeat -= 1
12
13    for _ in range(5):
14        print('{}hello {}'.format(threading.current_thread().name,datetime.now()))
15    print('{} leaves timer_func {}'.format(threading.current_thread().name,datetime.now()))
16
17 t = threading.Timer(2,timer_func2,(2,))
18 t.start()
19 for _ in range(20):
20     print('The main thread is doing something')
21     time.sleep(10)
```

```
The main thread is doing something
thread Thread-7 will start at 2018-11-09 14:35:20.674388
Thread-6,hello 2018-11-09 14:35:18.679388
Thread-6,hello 2018-11-09 14:35:18.679388
Thread-6,hello 2018-11-09 14:35:18.679388
Thread-6,hello 2018-11-09 14:35:18.679388
Thread-6,hello 2018-11-09 14:35:18.679388
Thread-6 leaves timer_func 2018-11-09 14:35:18.679388
thread Thread-8 will start at 2018-11-09 14:35:22.686503
Thread-7,hello 2018-11-09 14:35:20.691504
Thread-7,hello 2018-11-09 14:35:20.691504
Thread-7,hello 2018-11-09 14:35:20.691504
Thread-7,hello 2018-11-09 14:35:20.691504
Thread-7,hello 2018-11-09 14:35:20.691504
Thread-7 leaves timer_func 2018-11-09 14:35:20.691504
thread Thread-9 will start at 2018-11-09 14:35:24.698618
Thread-8,hello 2018-11-09 14:35:22.703619
Thread-8,hello 2018-11-09 14:35:22.703619
Thread-8,hello 2018-11-09 14:35:22.703619
Thread-8,hello 2018-11-09 14:35:22.703619
Thread-8,hello 2018-11-09 14:35:22.703619
Thread-8 leaves timer_func 2018-11-09 14:35:22.703619
thread Thread-10 will start at 2018-11-09 14:35:26.710733
Thread-9,hello 2018-11-09 14:35:24.716734
Thread-9,hello 2018-11-09 14:35:24.716734
Thread-9,hello 2018-11-09 14:35:24.716734
Thread-9,hello 2018-11-09 14:35:24.716734
Thread-9,hello 2018-11-09 14:35:24.716734
Thread-9 leaves timer_func 2018-11-09 14:35:24.716734
The main thread is doing something
thread Thread-11 will start at 2018-11-09 14:35:28.722848
Thread-10,hello 2018-11-09 14:35:26.727849
Thread-10,hello 2018-11-09 14:35:26.727849
Thread-10,hello 2018-11-09 14:35:26.727849
Thread-10,hello 2018-11-09 14:35:26.727849
Thread-10,hello 2018-11-09 14:35:26.727849
```

[illegible]

In [1]:

```
1 import threading
2 import time
3 from datetime import datetime, timedelta
4
5 def timer_func2(n,event_object):
6     while not event_object.isSet():
7         print('{}hello {}'.format(threading.current_thread().name,datetime.now()))
8         event_object.wait(n)
9         print('{} leaves timer_func {}'.format(threading.current_thread().name,datetime.now()))
10
11 event_object = threading.Event()
12 t1 = threading.Timer(2,timer_func2,(2,event_object))
13 t2 = threading.Timer(2,timer_func2,(2,event_object))
14 t1.start()
15 t2.start()
16 for _ in range(3):
17     print('The main thread is doing something')
18     time.sleep(5)
19 event_object.set()
20 t1.join()
21 t2.join()
```

```
The main thread is doing something
Thread-7,hello 2018-11-26 23:25:59.030114
Thread-6,hello 2018-11-26 23:25:59.030114
Thread-6,hello 2018-11-26 23:26:01.045145Thread-7,hello 2018-11-26 23:26:01.
045145
```

```
The main thread is doing something
Thread-7,hello 2018-11-26 23:26:03.060470Thread-6,hello 2018-11-26 23:26:03.
060470
```

```
Thread-6,hello 2018-11-26 23:26:05.060540Thread-7,hello 2018-11-26 23:26:05.
060540
```

```
The main thread is doing something
Thread-6,hello 2018-11-26 23:26:07.065305Thread-7,hello 2018-11-26 23:26:07.
065305
```

```
Thread-6,hello 2018-11-26 23:26:09.065971Thread-7,hello 2018-11-26 23:26:09.
065971
```

```
Thread-7,hello 2018-11-26 23:26:11.079962Thread-6,hello 2018-11-26 23:26:11.
079962
```

```
Thread-7 leaves timer_func 2018-11-26 23:26:12.045858Thread-6 leaves timer_f
unc 2018-11-26 23:26:12.045858
```

---

## multiprocessing.Process: create a process

```
import multiprocessing
aprocess = multiprocessing.Process(target=func,args(p1,...),kwargs={k1:v1,...})
```

- Call aprocess.start() to start aprocess.

- Call `aprocess.join(timeout)` to wait at most `timeout` seconds for the completion of `aprocess`.
- Use a lock to ensure that only one process is in the critical section.

`lock = multiprocessing.Lock()`

- Use queues and pipes to communicate between processes

<https://docs.python.org/3/library/multiprocessing.html#exchanging-objects-between-processes>

(<https://docs.python.org/3/library/multiprocessing.html#exchanging-objects-between-processes>)

## Pool Objects

**do not forget this line !!**

assign tasks to available processes automatically

```
import multiprocessing
import time
import numpy as np

def func(n,m):
    time.sleep(1)
    print("{} {}".format(multiprocessing.current_process().name))
    if n > m:
        n,m = m, n
    return sum(range(n,m))

if __name__ == "__main__":
    results = []
    all_data = [tuple(np.random.randint(0,100,(1,2)).ravel()) for _ in range(20)]
    with multiprocessing.Pool(multiprocessing.cpu_count()) as pool:
        returns = [pool.apply_async(func,args=(data[0],data[1])) for data in all_data]
        results = [a_return.get() for a_return in returns]
    print(results)
```

Get the return value for each task

SpawnPoolWorker-1  
 SpawnPoolWorker-6  
 SpawnPoolWorker-5  
 SpawnPoolWorker-8  
 SpawnPoolWorker-4  
 SpawnPoolWorker-3  
 SpawnPoolWorker-2  
 SpawnPoolWorker-7  
 SpawnPoolWorker-1  
 SpawnPoolWorker-6  
 SpawnPoolWorker-5  
 SpawnPoolWorker-8  
 SpawnPoolWorker-4  
 SpawnPoolWorker-3  
 SpawnPoolWorker-2  
 SpawnPoolWorker-7  
 SpawnPoolWorker-1  
 SpawnPoolWorker-8  
 SpawnPoolWorker-6  
 SpawnPoolWorker-5

## Executor Objects

- The `concurrent.futures` module provides a high-level interface for multithreading and multiprocessing.

<https://docs.python.org/3/library/concurrent.futures.html#module-concurrent.futures>

(<https://docs.python.org/3/library/concurrent.futures.html#module-concurrent.futures>)

- Executor Objects
  - `ThreadPoolExecutor(workers = number_of_workers)`
  - `ProcessPoolExecutor(workers = number_of_workers)`
- Methods for asynchronous execution
  - `submit`: the results will be yielded once they are completed.
  - `map`: the results are yielded in the same order of the data.

**submit**

**map**

**submit**

```
import concurrent.futures
import time

def echo(x):
    if x == 0:
        time.sleep(10)
    else:
        time.sleep(1)
    return x

begin = time.time()
with concurrent.futures.ThreadPoolExecutor(max_workers=10) as executor:
    fs = [executor.submit(echo, x) for x in range(100)]
    for future in concurrent.futures.as_completed(fs):
        print(future.result())
print(time.time()-begin)
```

this one is yielded first

18  
10  
26  
1  
5  
9  
13  
8  
...

futures will be yielded once they are completed.

ThreadPoolExecutor\_submit.py

**map**

```
import time
import concurrent.futures

def echo(x):
    if x == 0:
        time.sleep(10)
    else:
        time.sleep(1)
    return x

begin = time.time()
data = [x for x in range(100)]
with concurrent.futures.ThreadPoolExecutor(max_workers=100) as executor:
    #result = list(executor.map(echo, data)) # store the results in a list
    for i in executor.map(echo, data): # iterate over the result
        print(i)
print(time.time()-begin)
```

0  
1  
2  
3  
4  
5  
6  
7  
...

the results are yielded in the order consistent with data

ThreadPoolExecutor\_map.py

```
import concurrent.futures

def echo(x):
    if x == 0:
        time.sleep(10)
    else:
        time.sleep(1)
    return x
```

**submit**

```
if __name__ == '__main__':
    begin = time.time()
    with concurrent.futures.ProcessPoolExecutor(max_workers=10) as executor:
        fs = [executor.submit(echo, x) for x in range(100)]
        for future in concurrent.futures.as_completed(fs):
            print(future.result())
    print(time.time()-begin)
```

**map**

```
if __name__ == '__main__':
    begin = time.time()
    with concurrent.futures.ProcessPoolExecutor(max_workers=10) as executor:
        #result = list(executor.map(echo, data)) # store the results in a list
        for i in executor.map(echo, data): # iterate the result
            print(i)
    print(time.time()-begin)
```

## ProcessPoolExecutor vs ThreadPoolExecutor

In [2]:

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 def draw_result(title, workers, result):
4     relative_computing_time = result.copy()
5     starting_time = min(relative_computing_time[:,0])
6     relative_computing_time = relative_computing_time - starting_time
7     plt.barh(y=list(range(1,1+relative_computing_time.shape[0])),width=relative_computing_time[:,1])
8     plt.xticks([(i+1)//2 + sum(workers[0:idx]) for idx,i in enumerate(workers)], [str(i) for i in range(1,1+relative_computing_time.shape[0])])
9     plt.title(title)
10    plt.xlabel('seconds')
11    plt.ylabel('# of threads/processes')
12    plt.grid(True)
13    print('average computing time {}'.format(np.mean((relative_computing_time[:,1]-relative_computing_time[:,0])/(relative_computing_time[:,1]-relative_computing_time[:,0])))
14    plt.show()
```

ProcessPoolExecutor works on Ubuntu, but fails when running Jupyter notebook on Windows

<https://stackoverflow.com/questions/43836876/processpoolexecutor-works-on-ubuntu-but-fails-with-brokenprocesspool-when-r> (<https://stackoverflow.com/questions/43836876/processpoolexecutor-works-on-ubuntu-but-fails-with-brokenprocesspool-when-r>)

In [3]:

```
1 import os
2 import time
3 import math
4 import threading
5 import multiprocessing
6 from concurrent.futures import as_completed, ProcessPoolExecutor, ThreadPoolExecutor
7 import urllib.request
8 import winprocess
9
10 def download(x):
11     a = time.time()
12     with urllib.request.urlopen('http://python.org/') as response:
13         html = response.read()
14     b = time.time()
15     return [a,b]
16
17 def computing(x):
18     a = time.time()
19     x = 0
20     y = 0.001
21     while x < 5000000:
22         y = math.cos(y)
23         x = x + 1
24     b = time.time()
25     return [a,b]
```

## CPU bound tasks

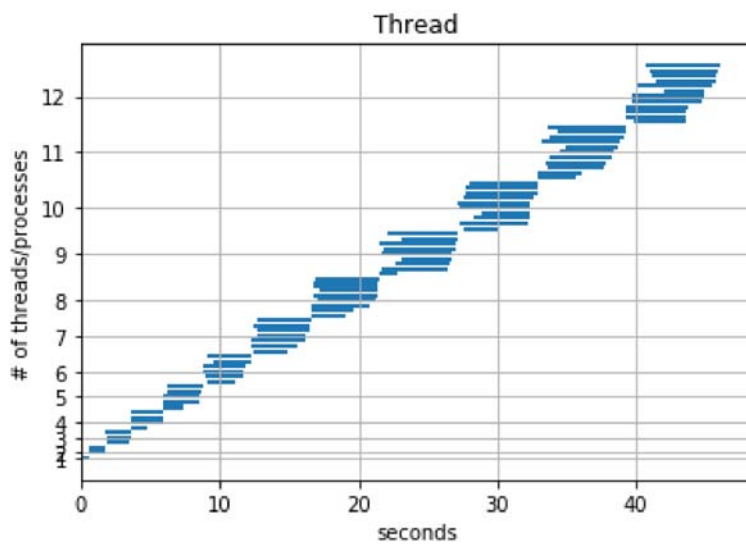
- The average computing time of the multiple threads is 3.68 seconds.
- The average computing time of the multiple processes is 0.92 seconds.



In [4]:

```
1 result = []
2 workers_lst = list(range(1,os.cpu_count()+1))
3 for workers in workers_lst:
4     with ThreadPoolExecutor(max_workers=workers) as executor:
5         futures = set()
6         for idx in range(workers):
7             future = winprocess.submit(
8                 executor, computing, idx
9             )
10            futures.add(future)
11
12        for future in as_completed(futures):
13            result.append(future.result())
14
15 draw_result('Thread',workers_lst,np.array(result))
```


average computing time 3.6844754830384865



In [5]:

```
1 result = []
2 workers_lst = list(range(1,os.cpu_count()+1))
3 for workers in workers_lst:
4     with ProcessPoolExecutor(max_workers=os.cpu_count()) as executor:
5         futures = set()
6         for idx in range(workers):
7             future = winprocess.submit(
8                 executor, computing, idx
9             )
10            futures.add(future)
11
12        for future in as_completed(futures):
13            result.append(future.result())
14
15 draw_result('Process',workers_lst,np.array(result))
```

average computing time 0.9212198165746835



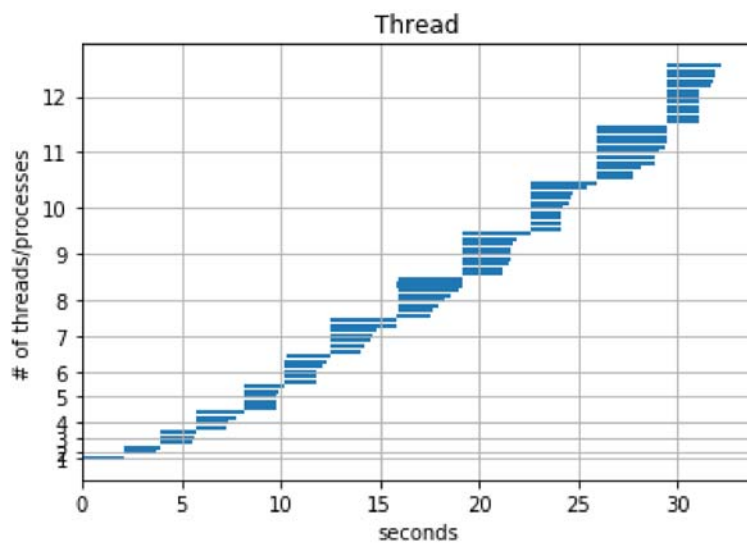
## I/O bound tasks

- The average computing time of the multiple threads is 2.21 seconds.
- The average computing time of the multiple processes is 2.39 seconds.

In [6]:

```
1 result = []
2 workers_lst = list(range(1,os.cpu_count()+1))
3 for workers in workers_lst:
4     with ThreadPoolExecutor(max_workers=workers) as executor:
5         futures = set()
6         for idx in range(workers):
7             future = winprocess.submit(
8                 executor, download, idx
9             )
10            futures.add(future)
11
12        for future in as_completed(futures):
13            result.append(future.result())
14
15 draw_result('Thread',workers_lst,np.array(result))
```

average computing time 2.2065395544736814



In [7]:

```
1 result = []
2 workers_lst = list(range(1,os.cpu_count()+1))
3 for workers in workers_lst:
4     with ProcessPoolExecutor(max_workers=os.cpu_count()) as executor:
5         futures = set()
6         for idx in range(workers):
7             future = winprocess.submit(
8                 executor, download, idx
9             )
10            futures.add(future)
11
12        for future in as_completed(futures):
13            result.append(future.result())
14
15 draw_result('Process',workers_lst,np.array(result))
```

average computing time 2.393884866665571

